

Generative Models: Clustering

Machine Learning

Unsupervised Learning I


- ▶ So far, we have looked at constructing models from data D of the form $(x_1, y_1), (x_2, y_2), \dots$ where the y_i are either numeric or nominal values, and the x_i are usually some vector of values
- ▶ The y_i are sometimes called *dependent* or *outcome* variables
- ▶ When the y_i are nominal (classes, for example), then the x_i 's with the same y value can be thought of as a *cluster* or a group
 - ▶ When we constructed a set of rules for predicting a class, we were identifying patterns that characterise that cluster
 - ▶ As we saw with classification models, the task of predicting y_i given the x_i can be seen as an application of discriminatory models (that is, we build a model for $P(Y|X)$)
 - ▶ We saw applications of this with logistic regression and class-probability trees
- ▶ But what if there were no y_i 's, and we still want to find groups of similar instances

Unsupervised Learning II

- ▶ One way to do this is to use generative models (that is, build a model for $P(X)$)
- ▶ We will see two examples of this approach using Naive Bayes and a mixture of Gaussians
- ▶ All instances in a cluster can then be taken as having the same class value
- ▶ Classification with missing (x_i, y_i) instances in which all the y_i are missing (cue: EM)
- ▶ BUT: why are there no y_i 's?
 - ▶ Finding labels can be expensive: easier to get large numbers of x_i 's
 - ▶ Even if some of the data were labelled, there might be *concept drift*
 - ▶ There may be no specific concept
- ▶ There are many different kinds of clustering models, but they are all essentially based on on a notion of *distance* between instances

Unsupervised Learning III

- Here are some different kinds of cluster representations:



From: Witten, Frank and Hall, *Data Mining* (2011)

Clustering

- ▶ The goal of clustering is to find sub-groups of N elements (instances)
- ▶ The sub-grpups can be partitions, but not always. Elements from same group (*cluster*) should have high similarity, and elements from different clusters low similarity
- ▶ Homogeneity and separation not well-defined. In practice, these concepts usually rely on some notion of distance between instances

Step Back: Distances I

- ▶ Minkowski distance: For instances from \mathbb{R}^d , the *Minkowski distance* of order $p > 0$ is defined as:

$$Dis_p(x, y) = \left(\sum_{j=1}^d |x_j - y_j|^p \right)^{1/p}$$

- ▶ This is sometimes denoted L_p . We will often refer to Dis_p simply as the p -norm
- ▶ The 2-norm refers to the familiar *Euclidean distance*:

$$Dis_2(x, y) = \sqrt{\sum_{j=1}^d (x_j - y_j)^2} = \sqrt{(x - y)^T (x - y)}$$

Step Back: Distances II

- ▶ The 1-norm refers to the familiar *Manhattan distance*, (or *cityblock distance*):

$$Dis[1](x, y) = \sum_{j=1}^d |x_j - y_j|$$

- ▶ If we now let p grow larger, the distance is dominated by the largest coordinate-wise distance. $Dis[\infty](x, y) = \max_j |x_j - y_j|$ this called the *Chebyshev distance*.
- ▶ You will sometimes see references to the 0-norm, which counts the number of non-zero elements in a vector. The corresponding distance then counts the number of positions in which vectors x and y differ

Step Back: Distances III

- ▶ This is not strictly a Minkowski distance; however, we can define it as

$$Dis[0](x, y) = \sum_{j=1}^d (x_j - y_j)^0 = \sum_{j=1}^d I(x_j \neq y_j)$$

where $I(.)$ is an indicator function, that is 1 if its argument is and 0 otherwise

- ▶ If x and y are binary strings, this is also called the *Hamming distance*.
- ▶ Alternatively, we can see the Hamming distance as the number of bits that need to be flipped to change x into y .
- ▶ For non-binary strings of unequal length this can be generalised to the notion of *edit distance* or the *Levenshtein distance*
- ▶ Some desirable properties of distances
 - ▶ Distance between two points cannot be negative

Step Back: Distances IV

- ▶ Distance between a point to itself is 0 (conversely, if the distance between a pair of points is 0, then the points are identical)
- ▶ Distance between points A and B is the same as between points B and A
- ▶ Distance between points A and C is less than or equal to the sum of the distances between points A and B; and B and C
- ▶ Distances that satisfy these properties are called *metrics*
- ▶ Given a set of points and a distance function, we can construct a *distance matrix* whose entry d_{ij} is the distance between x_i and x_j

Simple Distance-Based Clustering: Nearest Neighbour I

- ▶ The simplest kind of learning consists of looking up the table of instances seen so far:
 1. Given: A dataset of instances $\{x_1, x_2, \dots, x_k\}$ where each x_i is actually a N -dimensional vector (in which one of the attributes may be the dependent or outcome variable y_i)
 2. Find: the instance x^* that most closely resembles a new instance x_{new}
 3. Update: x_{new} is added to the set of stored instances
- ▶ THEN WHAT? Depends: if you are interested in classifying x_{new} , or in numeric prediction then return y_i
- ▶ What if there is more than one x_{new} is reasonably similar to several x_i 's? Return the (weighted) mean, mode, or median — whichever makes sense
- ▶ BUT: is this “learning”? Yes, in a manner (by memorisation, and using a distance measure)
 - ▶ Usually uses k closest instances (k -nearest-neighbour)

Simple Distance-Based Clustering: Nearest Neighbour II


- ▶ Instance-based learning
 - ▶ Storage: may end up storing a lot (all?) the instances
 - ▶ Distance: Standard Euclidean distance is OK if all attributes are numeric. But what if they are not all numeric?
 - ▶ Pattern: no explicit pattern is learnt of similar instances, so we will always need all of the data
- ▶ Clustering methods are a natural extension of instance-based learning, but using more powerful representations for groups of instances

Probabilistic Clustering: Simple Bayes Model I

- ▶ Distance-based measures can be seen as special-cases or approximations of a more general form of probability-based clustering it
 - ▶ Assume some joint-probability model
 - ▶ Estimate the parameters of the model that result in highest posterior probability (or highest likelihood, if using equi-probable priors)
- ▶ Before we look at more elaborate probabilistic models for clustering, we will look at a simple form of probabilistic clustering, using Naive Bayes as the joint probability model, that uses ideas from a method known as the EM algorithm

Probabilistic Clustering: Simple Bayes Model II

- Recall the network structure for the simple Bayes classifier:



- Usually, we are concerned with estimating parameters with this network given data for the random variables, along with a class value of this kind:

Probabilistic Clustering: Simple Bayes Model III

X_1	X_2	...	X_n	Y
1	0	...	0	+
1	1	...	0	-
0	1	...	1	-
...
...

We can easily compute the CPTs with this data (with beta priors if needed)

- ▶ We can treat the clustering problem as one for which the Y values are missing
 - ▶ How do we calculate the CPT entries now?
 - ▶ EM
- ▶ Let us look at an example with $N = 4$. Suppose the data are:

Probabilistic Clustering: Simple Bayes Model IV

X_1	X_2	X_3	X_4	Y
1	0	1	1	?
0	1	1	0	?
0	0	1	1	?
...

- ▶ The parameters that are needed for the Simple Bayes model are $P(Y)$, $P(X_1|Y)$, $P(X_2|Y)$, $P(X_3|Y)$ and $P(X_4|Y)$
- ▶ As usual with EM, we will augment the data with a *count* column:

X_1	X_2	X_3	X_4	Y	<i>Count</i>
1	0	1	1	?	
0	1	1	0	?	
0	0	1	1	?	
...	

Probabilistic Clustering: Simple Bayes Model V

- ▶ Each cluster will act as a coin. We do not know which coin was used to obtain each data instance
- ▶ SO: Each data instance has a separate row for each coin, with initial counts assigned randomly:

X_1	X_2	X_3	X_4	Y	Count
1	0	1	1	1	0.4
1	0	1	1	2	0.1
1	0	1	1	3	0.5
0	1	1	0	1	0.2
0	1	1	0	2	0.4
0	1	1	0	3	0.4
0	0	1	1	1	0.3
0	0	1	1	2	0.3
0	0	1	1	3	0.4
...

Probabilistic Clustering: Simple Bayes Model VI

NOTE: The counts are intended to be “fractional instances”. The total number of instances should sum to the number of instances of the combination. We will achieve this by rescaling appropriately

- ▶ In the E-step, counts are updated. In the M-step, parameters are obtained using Maximum Likelihood.
- ▶ Our model stores values for $P(C)$, $P(X_1|C)$, $P(X_2|C)$, $P(X_3|C)$, $P(X_4|C)$. Assume any particular combination r of X values happens m_r times

Probabilistic Clustering: Simple Bayes Model VII

Here is an algorithm sketch:

1. Assign $P(C), P(X_j|C)$ randomly
2. while not done
 - 2.1 For each combination r of $X_{1,r}, X_{2,r}, X_{3,r}, X_{4,r}$ values:
 - 2.1.1 Calculate $P(C = c_k | X_{1,r}, X_{2,r}, X_{3,r}, X_{4,r})$ using usual NB inference with current values of $P(C), P(X_j|C)$ Call this $P(c_k | X_r)$
 - 2.1.2 Update fractional counts.
 - 2.2 Re-estimate $P(C), P(X_j|C)$ using new counts

The EM Algorithm: Coins with Missing Data I

- ▶ Suppose we have data from a sequence of Bernoulli trials, each with a probability of success p . Previously, we have seen how the maximum likelihood estimate of p was s/n where s was the number of successes observed in the data, and n were the total number of trials
- ▶ Each Bernoulli trial is like tossing a biased coin with probability p of landing *heads*. Maximum likelihood estimation can be used with more than 1 coin. If we have two coins A and B , then to obtain the maximum likelihood estimates of the parameters p_A and p_B , we repeatedly do the following:
 1. Randomly pick a coin (for the moment with equal probability)
 2. Toss the coin some number (say 10) times
 3. Record the experiment number, the number of *heads* observed; and which coin was chosen
 4. Repeat the experiment (that is, return to Step 1)

The EM Algorithm: Coins with Missing Data II

The maximum likelihood estimates of p_A and p_B will simply be the proportion of *heads* on tosses for which A and B were used

- ▶ Now, consider a harder problem. You still have 2 coins A and B , and you have data from say 5 repetitions:

R1. HTTTHHTHTH

R2. HHHHTHHHHH

R3. HTHHHHHHTHH

R4. HTHTTTTHHTT

R5. THHHHTHHHTH

But there is no record of which coin was used in each case. How can we estimate p_A and p_B ? (This is from *Nature Biotechnology*, Vol 26, No.8, pp 897–899.)

- ▶ One way is the following:
 1. Start with some guess about p_A and p_B . Call this $p_A^{(0)}$ and $p_B^{(0)}$

The EM Algorithm: Coins with Missing Data III

2. For each of the repeats R1–R5, calculate $P_A = P(D|p_A^{(0)})$ and $P_B = P(D|p_B^{(0)})$ (where D is the sequence of H 's and T 's on that repetition). If $P_A > P_B$, then assume A was used, otherwise assume B was used.
3. Now we have a complete table, and can calculate maximum likelihood estimates as before

We could iterate this entire procedure, using the maximum likelihood estimate as the guesses in Step (1).

- ▶ The “expectation-maximisation” algorithm or EM algorithm is a refinement of this basic idea
- ▶ The current guesses for p_A and p_B are used to obtain new “weighted” instances.

The EM Algorithm: Coins with Missing Data IV

- ▶ The weights are proportional to $P(\text{coin} = A|D)$ and $P(\text{coin} = B|D)$. These are proportional to $P(D|\text{coin} = A)P(\text{coin} = A)$ and $P(D|\text{coin} = B)P(\text{coin} = B)$. Since the coins were known to be chosen with equal probability, then $P(\text{coin} = A|D) = \alpha P(D|\text{coin} = A)$ and $P(\text{coin} = B|D) = \alpha P(D|\text{coin} = B)$. We can calculate this using the current guess of p_A and p_B .
- ▶ For example, with an initial guess of 0.60 for p_A and 0.50 for p_B , we find the following
 - R1. $D1 = HTTTHHTH$. Then, $P(\text{Coin} = A|D1) = 0.45$ and $P(\text{Coin} = B|D1) = 0.55$. So, we create a new set of weighted instances generated by both coins:

CoinA : 2.2*H* and 2.2*T*

CoinB : 2.8*H* and 2.8*T*

(Both):] 5*H* and 5*T* (as before)

The EM Algorithm: Coins with Missing Data V

R2. $D2 = HHHHTHHHHH$ Then, $P(\text{Coin} = A|D2) = 0.80$ and $P(\text{Coin} = B|D2) = 0.2$, and get weighted instances:

$\text{Coin}A : 7.2H$ and $0.8T$

$\text{Coin}B : 1.8H$ and $0.2T$

(Both): $9H$ and $1T$ (as before)

and so on, for repetitions R3–R5. In each case, the number of *heads* in the weighted instances is the expected value of the number of *heads*, given the current estimate of the parameters. The totals obtained after all calculations are:

$\text{Coin}A : 21.3H$ and $8.6T$

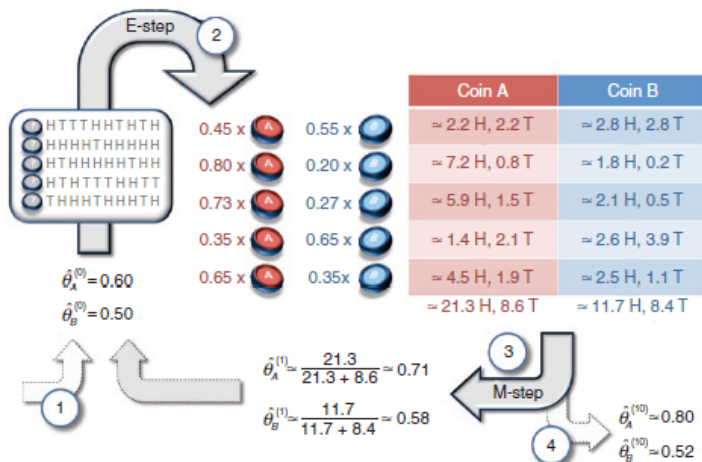
$\text{Coin}B : 11.7H$ and $8.4T$

This gives a new estimate of p_A as $21.3/(21.3 + 8.6) = 0.71$ and p_B as 0.58

The EM Algorithm: Coins with Missing Data VI

- ▶ This can be used iteratively. The procedure we have just done is one iteration of the EM algorithm

The EM Algorithm: Coins with Missing Data VII



Probabilistic Clustering: Mixture Models I

- ▶ Assume we have a model structure with some parameter Θ , and we want to work out which value of $\Theta = \theta$ is best for the data x (previously we used d) we have observed.
 - ▶ As before, θ and x will be values of random variables Θ and X (used to be D)
 - ▶ Both Θ and X are now continuous r.v.'s
- ▶ The only thing that changes is that both likelihood and priors are calculated using p.d.f's. That is:

Hypothesis. $\Theta = \theta$

Data. X in an interval dx around x

Prior. $f(\theta)d\theta$

Likelihood. $f(x|\theta)dx$ (correctly, $f_{X|\Theta}(x|\theta)dx$)

Numerator. $f(x|\theta)dx f(\theta)d\theta$

Denominator. $\int_0^1 f(x|\theta)dx f(\theta)d\theta$

Posterior. $f(\theta|x)d\theta$

Probabilistic Clustering: Mixture Models II

► Example

► Suppose $X \sim N(\theta, 1)$ and $\Theta \sim N(2, 1)$

► Then:

Prior. $c_1 e^{-(\theta-2)^2/2} d\theta$

Likelihood. $c_2 e^{-(5-\theta)^2/2} dx$

Numerator. $c_3 e^{-(2\theta^2-14\theta+29)/2} dx d\theta =$
 $c_3 e^{-((\theta^2-7/2)^2+9/4)} dx d\theta = c_4 e^{-(\theta^2-7/2)^2} dx d\theta$

Denominator. $(\int c_4 e^{-(\theta^2-7/2)^2} d\theta) dx = c_5 dx$

Posterior. $c_6 e^{-(\theta^2-7/2)^2} d\theta$


► The posterior p.d.f. is $N(7/2, \sigma)$ where $2\sigma^2 = 1$ and $c_6 = 1/(\sigma\sqrt{2\pi})$

► So a Gaussian is conjugate prior for Gaussian likelihood.

Mixture Models I

- ▶ A *mixture* is a set of k probability distributions, each representing a cluster
- ▶ Any particular instance x_i belongs to a cluster, but it is not known which one. So, each distribution j is used to compute the probability that x_i belongs to cluster j
- ▶ Simplest mixture model: there is one numeric attribute X and 2 or more clusters, each specified by a Gaussian with different means and variances

Mixture Models II



(From: Witten, Frank and Hall, (2012))

- ▶ Given a dataset, if we knew which cluster each instance came from, then estimating the means and s.d's of the Gaussians is easy:

$$\mu_j = \sum_{x_i \in C_j} \frac{x_i}{c_j} \quad \sigma_j^2 = \sum_{x_i \in C_j} \frac{(x_i - \mu_j)^2}{c_j - 1}$$

Mixture Models III

- ▶ Once we have the estimates of the parameters, finding the cluster for an instance x is also easy:

$$P(j|x) = \frac{P(x|j)P(j)}{P(x)} = \frac{N(x; \mu_j, \sigma_j)P(j)}{P(x)}$$

where:

$$N(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{\frac{-(x-\mu)^2}{2\sigma^2}}$$

(strictly speaking, this is not correct of course, since the $N(.,.)$ is a p.d.f. But if we put in the $d\theta$'s, this is OK)

- ▶ BUT: we don't know the cluster to which instance belongs. So, what is to be done (CUE: EM)
 - ▶ Start with guesses for the μ_j, σ_j and $P(j)$ (for 2 clusters, only one guess is needed for $P(j)$).

Mixture Models IV

- ▶ Then calculate the fractional instances belonging to each cluster (these will be taken to be equal to the probability of instances belonging to a class)
- ▶ Re-calculate the means and s.d's:

$$\mu_j = \frac{\sum_{i=1}^N w_i x_i}{\sum_i w_i}$$

$$\sigma_j^2 = \frac{\sum_{i=1}^N w_i (x_i - \mu_j)^2}{\sum_i w_i}$$

- ▶ The “best clustering” is one in which all instances are in the correct cluster. That is $P(j|x) = 1$ if $x \in j$. The value of $P(x|j)P(j)$ will be a maximum in this case
- ▶ The procedure iterates using a measure of goodness based on the weighted likelihood L of the data. The weight of cluster j is the weight of instances belonging to cluster j . The weight of an instance x_i in cluster j is taken as $P(j)P(x_i|j)$.

Mixture Models V

- ▶ Again, if we use $N(.,.)$ for $P(x_i|j)$, this is technically not correct, and L will not be between 0 and 1. But it is a good enough approximation for directing the EM procedure
 - ▶ In practice, $\log L$ will be used to avoid floating point errors
 - ▶ Typically, the EM procedure will run for about 10 iterations or so: usually $\log L$ will increase sharply and then subsequent changes will be small
- ▶ Extending from 2 to K classes is straightforward if we know K in advance
- ▶ If there are multiple numeric attributes, then each instance is really a vector x_i and either the usual Naive Bayes assumption is made:

$$P(x_i|j) = \prod_{d=1}^D P(x_d|j)$$

OR the clusters are treated as mixtures of multivariate Normal distributions. This allows the calculation of joint-probabilities of a set of correlated attributes.

- ▶ If there are missing values, then either impute using means, modes, medians, or add more machinery to the EM procedure
- ▶ If the Gaussian is inappropriate, use some other p.d.f (for example, a sigmoid function)
- ▶ Increasing the covariant attributes increases the number of parameters (for every pair, the σ_{ij} is replaced by Σ_{ij} which has 4 entries rather than 1)
 - ▶ More parameters means greater risk of overfitting. More correlated attributes mean more parameters.

Approximate Mixture Modelling: K-Means I

- ▶ An approximation to the full EM-based Gaussian mixture modelling
 - ▶ Single parameter: means of clusters
 - ▶ Distance-based cluster membership rather than p.d.f-based cluster membership
 - ▶ Instance belongs to a single cluster, rather than fractional instances
- ▶ Procedure:
 1. Set value for K (the number of clusters)
 2. Randomly assign values for the centres (means) of the k clusters
 3. repeat
 4. For each instance x do:
 - 4.1 Calculate distance to each of the k -means
 - 4.2 Assign x to the closest cluster
 5. Recalculate the arithmetic means of each of the k -clusters
 6. until (local minimum)

Approximate Mixture Modelling: K-Means II

- ▶ BUT: when are we done? The procedure stops at a local minimum for:

$$D = \sum_{n=1}^N \sum_{k=1}^K z_{nk} \|x_n - \mu_k\|^2$$

where $z_{nk} = 1$ if x_n is closest to cluster k and 0 otherwise; and $\|\cdot\|$ is the 2-norm

- ▶ Given a set of points, the arithmetic mean of the points minimises the sum of squared Euclidean distances to those points (Can you prove this?)
- ▶ Given a set of points C_k in a cluster k with a centre d_k , $\sum_{j \in C_k} \|x_j - d_k\|^2$ will be minimum iff d_k is the arithmetic mean of the points in C_k


Approximate Mixture Modelling: K-Means III

- ▶ When the cluster assignment for each point does not change after an iteration (that is, if z_{nk} does not change), then by construction d_k will be the arithmetic mean of the points in C_k
- ▶ BUT: as the number of clusters increases, D becomes smaller (why?) In the worst case, it is possible to end up with $D = 0$ for $K = N$ (clusters with single points)
 - ▶ This is a problem

More Problems: When K -means Can Fail I

- ▶ K -means has some underlying assumptions:
 - ▶ Clusters are spheres centred at different points
 - ▶ Clusters are approximately of the same size (priors are the same)
- ▶ Each of these assumptions can be violated quite easily. Here is a case where we know the number of clusters exactly, and even their arithmetic means :


More Problems: When K -means Can Fail II



(from stats.stackexchange.com)


More Problems: When K -means Can Fail III

But the data do not satisfy the “spherical data” assumption, and K -means fails:




More Problems: When K -means Can Fail IV

- Here there are unequal clusters:



More Problems: When K -means Can Fail V

and again, the procedure finds sub-optimal clusters




- SO: what is to be done? There are other clustering techniques (that make other assumptions)

Hierarchical Clustering I

- ▶ Assign each instance to its own cluster. There are now N clusters
- ▶ Find the closest pair of clusters and merge them into a single cluster
 - ▶ Distances between clusters need not just be between the arithmetic mean of instances in the cluster
 - ▶ Each merge operation reduces the clusters by 1
- ▶ Distance between clusters
 - ▶ Single-linkage: nearest neighbour
 - ▶ Complete-linkage: furthest neighbour
 - ▶ Average-linkage: average of distances of all instance-pairs
- ▶ SO: pick a distance function, pick a linkage method, find distances between all cluster pairs, and merge the closest pair
 - ▶ This is *agglomerative* clustering (goes from N clusters to 1 cluster). The complexity is $O(N^2 \log N)$

Hierarchical Clustering II

- ▶ The other way is *divisive* clustering (goes from 1 cluster to N clusters). The complexity is $O(2^N)$, although greedy methods can do better
- ▶ Single-linkage agglomerative clustering (using Euclidean distances)



Summary

- ▶ There are several problems which are not about predicting a value y , given a data instance x . Instead, the task is simply to find subsets S_1, S_2, \dots, S_k of a set of data instances $D = x_1, x_2, \dots, x_N$, s.t. the S_i satisfy some properties
- ▶ Problems like these are collectively called *unsupervised learning* tasks, of which *clustering* problems are the best known (but not the only one) They are best tackled by thinking of them as generative models rather than discriminative models
- ▶ At the heart of most generative models for clustering is a distance measure
- ▶ One easy way of looking at clustering is to consider it as an application of using generative probabilistic models (like Naive Bayes), with a probability-based distance measure. Other clustering techniques exist that identify subsets differently, using other kinds of distance measures (e.g. hierarchical clustering)